



Main pressure comparison · 12 matched workloads / 144 streams per policy

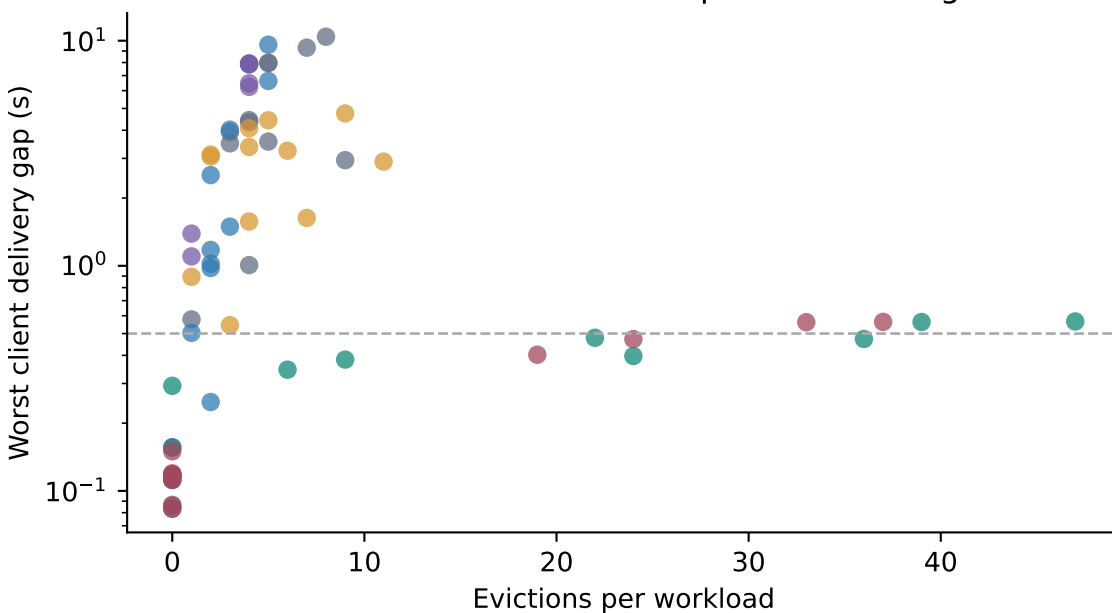
Do already-started streams keep getting service?



Dots: workloads · black mark: geometric mean



Eviction counts and interruption can diverge



The first-token waiting cost

